

# Red Hat OpenShift Container Platform

## On IBM Z and LinuxONE

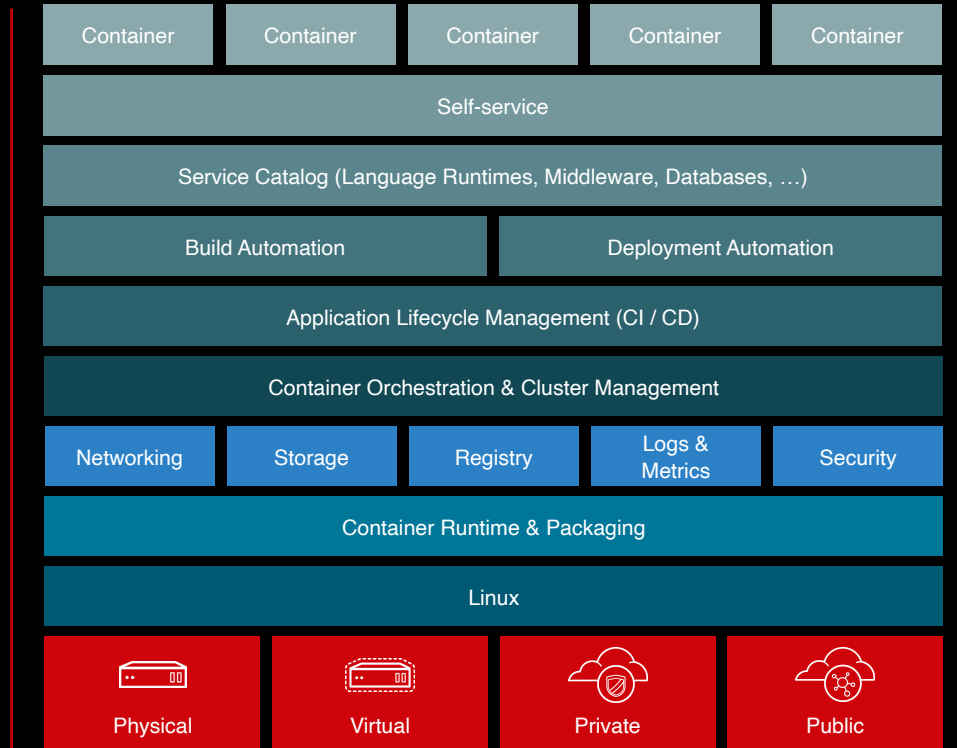
KVM Edition

---

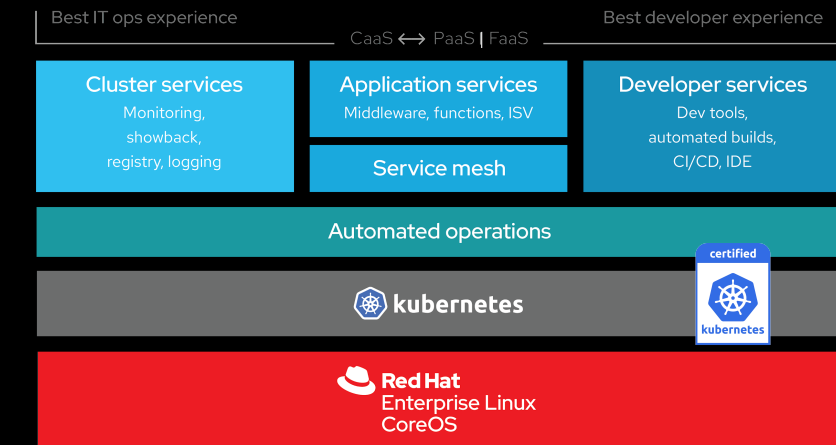
Filipe Miranda

Cloud Solutions Architect Leader, zAcceleration  
Red Hat Synergy  
fmiranda@ibm.com

# Clients are facing DIY challenges



# Red Hat OpenShift Available on IBM Z and LinuxONE




**Automated, full-stack installation** from the container host to application services

**Seamless Kubernetes deployment** to any cloud or on-premises environment

**Autoscaling** of cloud resources

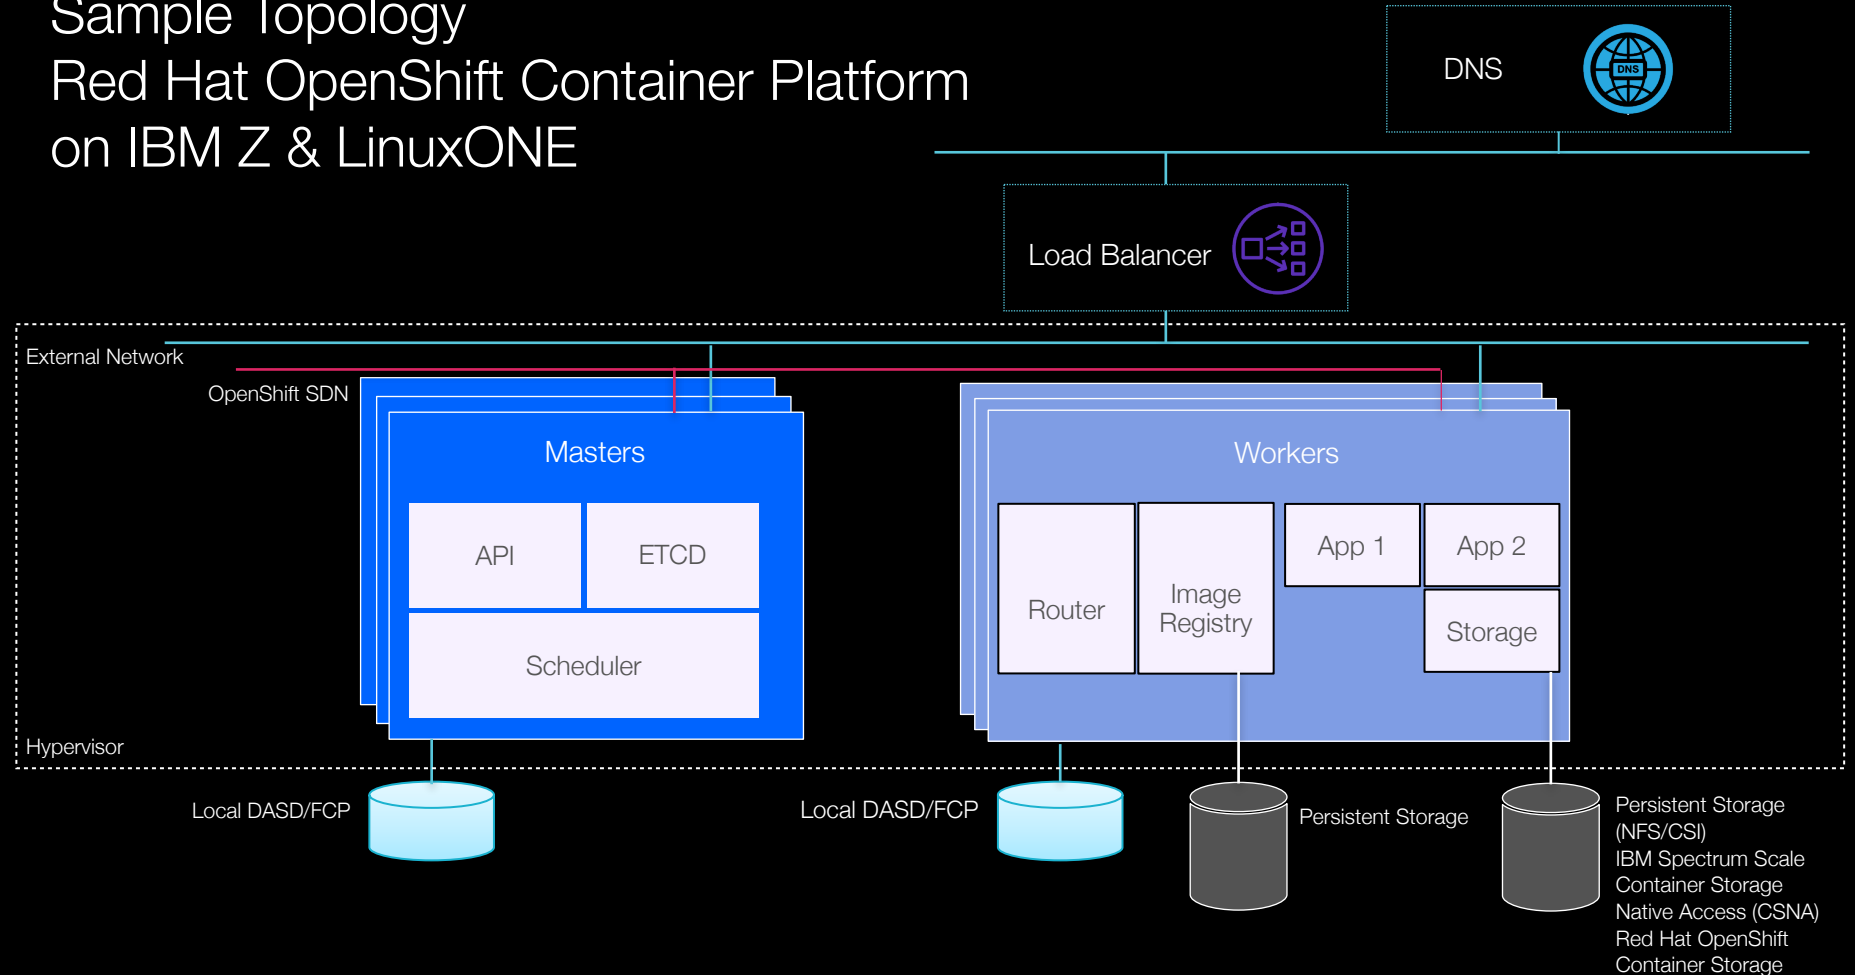
**One-click updates** for platform, services, and applications



# Sample Topology

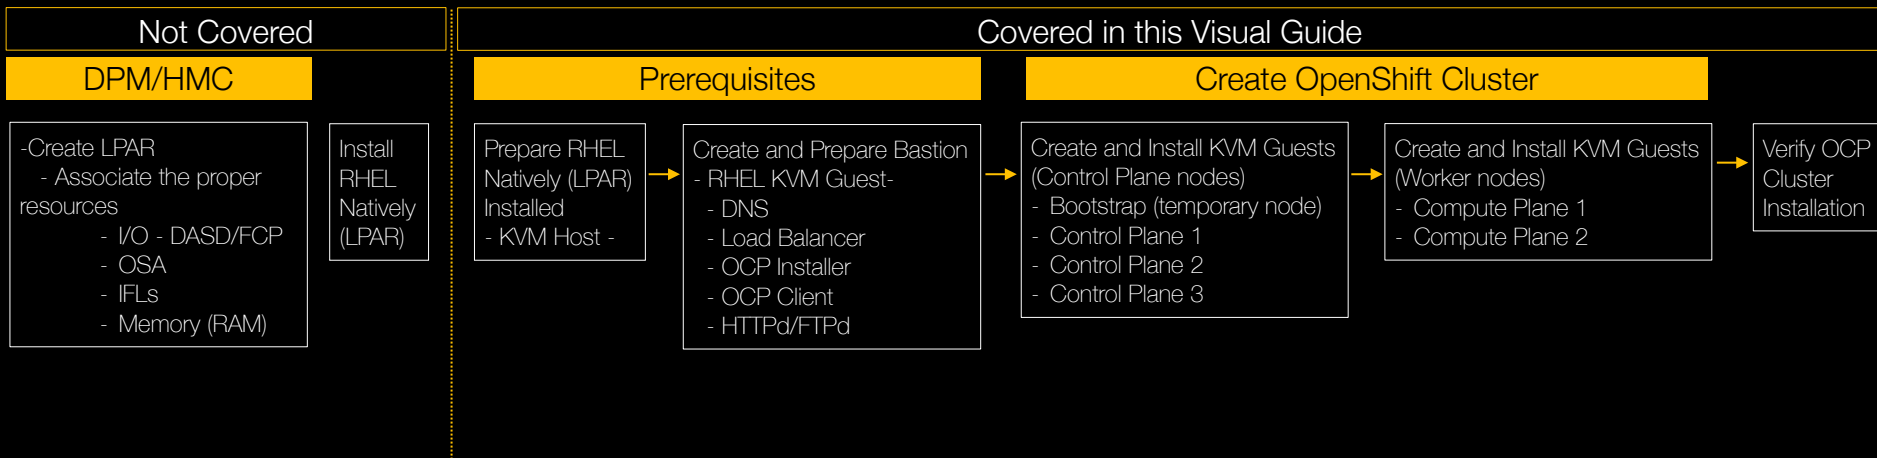
## Red Hat OpenShift Container Platform

### on IBM Z & LinuxONE




# Red Hat OpenShift Container Platform Install Process

KVM




# Red Hat OpenShift Container Platform Install Process

KVM



# KVM Support installation



- 1 Make sure you make these packages installed and that libvirt is started:

```
# yum install libvirt libvirt-devel libvirt-daemon-kvm qemu-kvm virt-manager libvirt-daemon-config-network  
libvirt-client qemu-img  
# systemctl enable --now libvirt  
# systemctl status libvirtd.service
```


- 2 

```
# systemctl status libvirtd
```

  - libvirtd.service - Virtualization daemon

```
Loaded: loaded (/usr/lib/systemd/system/libvirtd.service; enabled; vendor preset: enabled)  
Active: active (running) since Thu 2020-12-17 22:26:01 EST; 2 months 19 days ago  
Docs: man:libvirtd(8)  
      https://libvirt.org  
  
Main PID: 933606 (libvirtd)  
Tasks: 18 (limit: 32768)  
Memory: 68.9M  
CGroup: /system.slice/libvirtd.service  
└─933606 /usr/sbin/libvirtd --timeout 120
```

# Host KVM Network Preparation



1 Let's find details about our network interface:

```
...
3: enc4100: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
   link/ether 7e:4c:52:67:a6:f0 brd ff:ff:ff:ff:ff:ff
   inet <IP_ADDRESS>/24 brd <BROADCAST> scope global noprefixroute enc4100
       valid_lft forever preferred_lft forever

   inet6 fe80::7c4c:52ff:fe67:a6f0/64 scope link
       valid_lft forever preferred_lft forever
...
```

2 Create this xml file:


```
<network>
  <name>macvtap-net</name>
  <forward dev='enc4100' mode='bridge'>
    <interface dev='enc4100' />
  </forward>
</network>
```

3 Set up a bridge to act as a macvtap interface to the network:

```
# virsh net-create macvtap.xml
# virsh net-start --network macvtap-net
# virsh net-autostart --network macvtap-net
# virsh net-list --all
```



# Create and Configure Bastion



- 1 Download the RHEL ISO image to your RHEL KVM:

```
# wget /URL/rhel-8.3-s390x-dvd.iso
# mv rhel-8.3-s390x-dvd.iso rhel83.iso
```


- 2 Start the install process

```
virt# virt-install --connect qemu:///system --name bastion --memory 4096 --vcpus 2 --disk size=20 --cdrom /var/lib/libvirt/images/rhel83.iso
--accelerate --import --network network=macvtap-net --extra-args "ip=172.16.10.212::172.16.10.1:255.255.255.0:bastion.ocp.home.local::none
nameserver=172.16.10.38 vnc vncpassword=12341234 inst.repo=hd:/dev/vda ipv6.disable=1" --location /rhcos-install --qemu-commandline="-drive
if=none,id=ignition,format=raw,file=/var/lib/libvirt/images/rhel83.iso,readonly=on -device virtio-blk,serial=ignition,drive=ignition" --
noautoconsole
```

- 3 Use a local (laptop) VNC Viewer to connect to the Bastion VM to complete the RHEL Bastion install process.

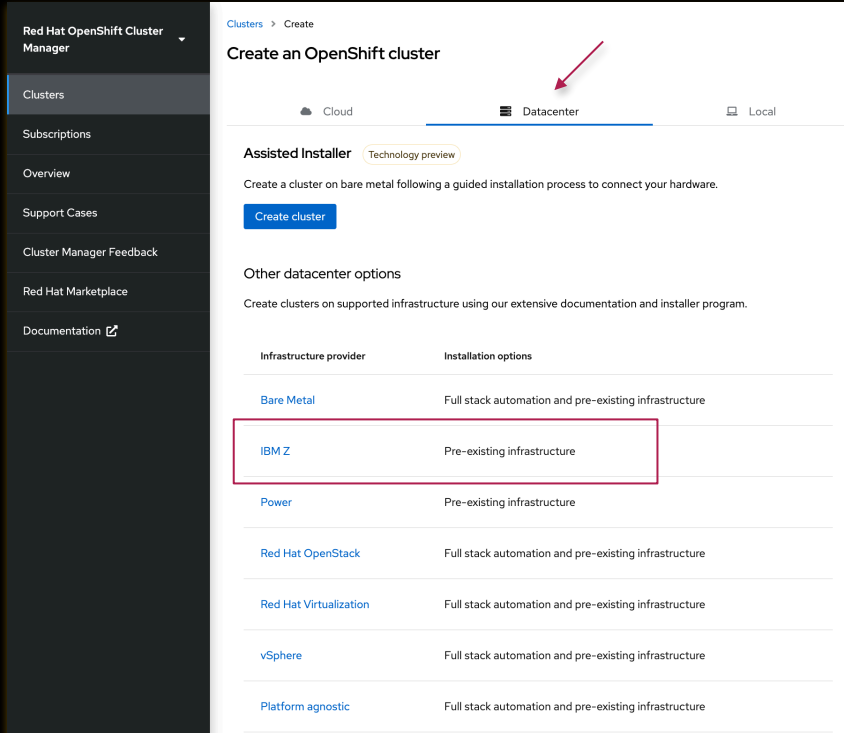
# Red Hat OpenShift Container Platform Install Process

KVM



# Download Software

<https://cloud.redhat.com/>



Red Hat OpenShift Cluster Manager

Clusters

Subscriptions

Overview

Support Cases

Cluster Manager Feedback

Red Hat Marketplace

Documentation

Clusters > Create

## Create an OpenShift cluster

Cloud Datacenter Local

Assisted Installer Technology preview

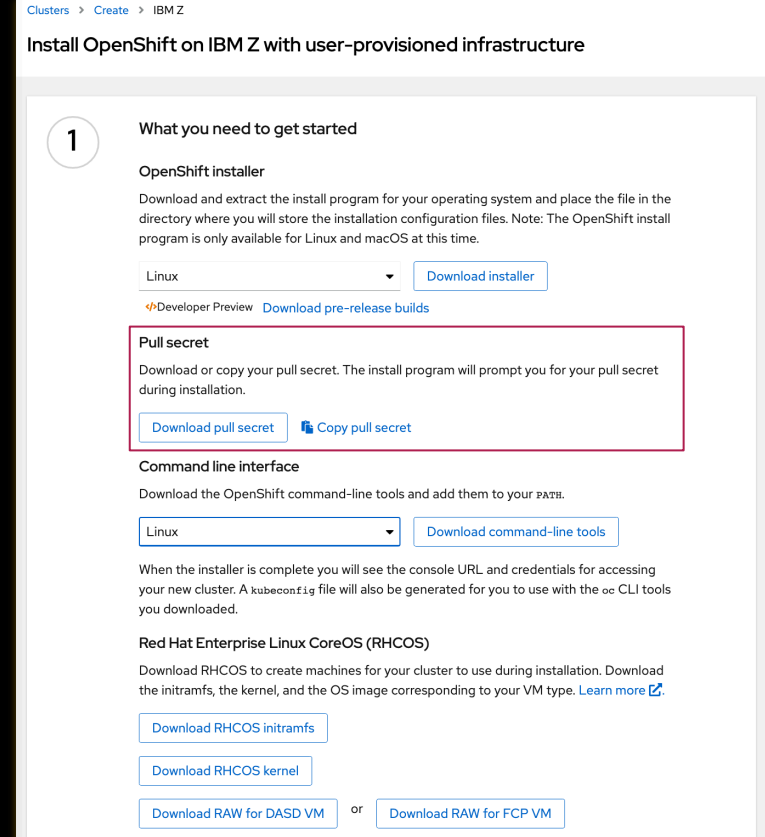
Create a cluster on bare metal following a guided installation process to connect your hardware.

Create cluster

### Other datacenter options

Create clusters on supported infrastructure using our extensive documentation and installer program.

Infrastructure provider	Installation options
Bare Metal	Full stack automation and pre-existing infrastructure
IBM Z	Pre-existing infrastructure
Power	Pre-existing infrastructure
Red Hat OpenStack	Full stack automation and pre-existing infrastructure
Red Hat Virtualization	Full stack automation and pre-existing infrastructure
vSphere	Full stack automation and pre-existing infrastructure
Platform agnostic	Full stack automation and pre-existing infrastructure



Clusters > Create > IBM Z

## Install OpenShift on IBM Z with user-provisioned infrastructure

### 1 What you need to get started

#### OpenShift installer

Download and extract the install program for your operating system and place the file in the directory where you will store the installation configuration files. Note: The OpenShift install program is only available for Linux and macOS at this time.

Linux Download installer

Developer Preview Download pre-release builds

#### Pull secret

Download or copy your pull secret. The install program will prompt you for your pull secret during installation.

Download pull secret Copy pull secret

#### Command line interface

Download the OpenShift command-line tools and add them to your PATH.

Linux Download command-line tools

When the installer is complete you will see the console URL and credentials for accessing your new cluster. A `kubeconfig` file will also be generated for you to use with the oc CLI tools you downloaded.

#### Red Hat Enterprise Linux CoreOS (RHCOS)

Download RHCOS to create machines for your cluster to use during installation. Download the initramfs, the kernel, and the OS image corresponding to your VM type. [Learn more](#)

Download RHCOS initramfs

Download RHCOS kernel


Download RAW for DASD VM or Download RAW for FCP VM

<https://mirror.openshift.com/pub/openshift-v4/s390x/clients/ocp/latest/>

<https://mirror.openshift.com/pub/openshift-v4/s390x/dependencies/rhcos/latest/latest/>

# DNS Requirements and Configuration Example:

Table 5. Required DNS records

Component	Record	Description
Kubernetes API	<code>api.&lt;cluster_name&gt;.&lt;base_domain&gt;.</code>	Add a DNS A/AAAA or CNAME record, and a DNS PTR record, to identify the load balancer for the control plane machines. These records must be resolvable by both clients external to the cluster and from all the nodes within the cluster.
	<code>api-int.&lt;cluster_name&gt;.&lt;base_domain&gt;.</code>	Add a DNS A/AAAA or CNAME record, and a DNS PTR record, to identify the load balancer for the control plane machines. These records must be resolvable from all the nodes within the cluster. <div>  <b>IMPORTANT</b>  The API server must be able to resolve the worker nodes by the host names that are recorded in Kubernetes. If the API server cannot resolve the node names, then proxied API calls can fail, and you cannot retrieve logs from pods. </div>
Routes	<code>*.apps.&lt;cluster_name&gt;.&lt;base_domain&gt;.</code>	Add a wildcard DNS A/AAAA or CNAME record that refers to the load balancer that targets the machines that run the Ingress router pods, which are the worker nodes by default. These records must be resolvable by both clients external to the cluster and from all the nodes within the cluster.
Bootstrap	<code>bootstrap.&lt;cluster_name&gt;.&lt;base_domain&gt;.</code>	Add a DNS A/AAAA or CNAME record, and a DNS PTR record, to identify the bootstrap machine. These records must be resolvable by the nodes within the cluster.
Master hosts	<code>&lt;master&gt;&lt;n&gt;.&lt;cluster_name&gt;.&lt;base_domain&gt;.</code>	Add DNS A/AAAA or CNAME records and DNS PTR records to identify each machine for the master nodes. These records must be resolvable by the nodes within the cluster.
Worker hosts	<code>&lt;worker&gt;&lt;n&gt;.&lt;cluster_name&gt;.&lt;base_domain&gt;.</code>	Add DNS A/AAAA or CNAME records and DNS PTR records to identify each machine for the worker nodes. These records must be resolvable by the nodes within the cluster.

```
@ IN SOA ns1.<domain>. admin.<domain>. (
    2020021821 ;Serial
    3600 ;Refresh
    1800 ;Retry
    604800 ;Expire
    86400 ;Minimum TTL
)

;Name Server Information
@ IN NS ns1.<domain>.

;IP Address for Name Server
ns1 IN A <DNS_server_IP_address>

; entry for the bootstrap host.
bootstrap.<cluster_name> IN A <bootstrap_IP_address>

; entry of your load balancer
haproxy IN A <loadbalancer_IP_address>

; entries for the master hosts
<control plane 1>.<cluster_name> IN A <control_plane1_IP_address>
<control plane 2>.<cluster_name> IN A <control_plane2_IP_address>
<control plane 3>.<cluster_name> IN A <control_plane3_IP_address>

; entry for the bastion host
bastion IN A <Infra_server_IP_address>

; entries for the workers hosts
<compute plane 1>.<cluster_name> IN A <compute_plane1_IP_address>
<compute plane 2>.<cluster_name> IN A <compute_plane2_IP_address>

; The api identifies the IP of your load balancer.
api.<cluster_name> IN CNAME haproxy.<domain>.
api-int.<cluster_name> IN CNAME haproxy.<domain>.

; The wildcard also identifies the load balancer.
*.apps.<cluster_name> IN CNAME haproxy.<domain>.
```

# Load Balancer

## HAProxy Example:

Table 3. API load balancer

Port	Back-end machines (pool members)	Internal	External	Description
6443	Bootstrap and control plane. You remove the bootstrap machine from the load balancer after the bootstrap machine initializes the cluster control plane. You must configure the <code>/readyz</code> endpoint for the API server health check probe.	X	X	Kubernetes API server
22623	Bootstrap and control plane. You remove the bootstrap machine from the load balancer after the bootstrap machine initializes the cluster control plane.	X		Machine config server

Table 4. Application Ingress load balancer

Port	Back-end machines (pool members)	Internal	External	Description
443	The machines that run the Ingress router pods, compute, or worker, by default.	X	X	HTTPS traffic
80	The machines that run the Ingress router pods, compute, or worker, by default.	X	X	HTTP traffic

```
1 # dnf install -y haproxy
```

```
2 /etc/haproxy/haproxy.cfg
```


```
listen ingress-http
  bind *:80
  mode tcp
  server worker0 <worker0_IP>:80 check
  server worker1 <worker1_IP>:80 check
```

```
listen ingress-https
  bind *:443
  mode tcp
  server worker0 <worker0_IP>:443 check
  server worker1 <worker1_IP>:443 check
```

```
listen api
  bind *:6443
  mode tcp
  server bootstrap <bootstrap_IP>:6443 check
  server master0 <master0_IP>:6443 check
  server master1 <master1_IP>:6443 check
  server master2 <master2_IP>:6443 check
```

```
listen api-int
  bind *:22623
  mode tcp
  server bootstrap <bootstrap_IP>:22623 check
  server master0 <master0_IP>:22623 check
  server master1 <master1_IP>:22623 check
  server master2 <master2_IP>:22623 check
```

# Create and configure the HTTP server



## NOTE

Always check the latest versions of OpenShift, the file names will change as new versions are made available:  
<https://mirror.openshift.com/pub/openshift-v4/s390x/dependencies/rhcos/latest/latest/>

1 `# dnf install -y httpd`

2 Change default port to 8080

3

```
# mkdir /var/www/html/bin /var/www/html/bootstrap


# wget https://mirror.openshift.com/pub/openshift-v4/s390x/dependencies/rhcos/latest/latest/rhcos-4.7.7-s390x-live-kernel-s390x-initramfs.s390x.img -O /var/www/html/bin/rhcos-kernel

# wget https://mirror.openshift.com/pub/openshift-v4/s390x/dependencies/rhcos/latest/latest/rhcos-4.7.7-s390x-live-initramfs.s390x.img -O /var/www/html/bin/rhcos-initramfs.img

# wget https://mirror.openshift.com/pub/openshift-v4/s390x/dependencies/rhcos/latest/latest/rhcos-4.7.7-s390x-live-rootfs.s390x.img -O rhcos-rootfs.img
```

4 `# systemctl enable --now httpd; systemctl status httpd`

# Installer and oc Client Tools




- 1 

```
# wget https://mirror.openshift.com/pub/openshift-v4/s390x/clients/ocp/latest/openshift-client-linux.tar.gz
tar -xvzf openshift-client-linux.tar.gz
```
- 2 

```
# wget https://mirror.openshift.com/pub/openshift-v4/s390x/clients/ocp/latest/openshift-install-linux.tar.gz
tar -xvzf openshift-install-linux.tar.gz
```
- 3 

```
# chmod +x kubect1 oc openshift-install
# mv kubect1 oc openshift-install /usr/local/bin/
```

# Install-config.yaml



## Pull secret

Download or copy your pull secret. The install program will prompt you for your pull secret during installation.

[Download pull secret](#)[Copy pull secret](#)

If you do not have an SSH key that is configured for password-less authentication on your computer, create one. For example, on a computer that uses a Linux operating system, run the following command:

```
# ssh-keygen -t rsa -b 4096 -N ''
```

## install-config.yaml

```
apiVersion: v1
baseDomain: <domain>
compute:
- architecture: s390x
  hyperthreading: Enabled
  name: worker
  replicas: 0
controlPlane:
  architecture: s390x
  hyperthreading: Enabled
  name: master
  replicas: 3
metadata:
  name: <cluster_name>
networking:
  clusterNetwork:
  - cidr: 10.128.0.0/14
    hostPrefix: 23
  networkType: OpenShiftSDN
  serviceNetwork:
  - 172.30.0.0/16
platform:
  none: {}
fips: false
pullSecret: '<pull-secret>'
sshKey: '<ssh-public-key>'
```



# Install-config.yaml



## Prerequisites

Prepare RHEL  
Natively (LPAR)  
Installed  
- KVM Host


Create and Prepare Bastion  
- RHEL KVM Guest  
- DNS  
- Load Balancer  
- OCP Installer  
- OCP Client  
- HTTPd/FTPd

## Create OpenShift Cluster

Create and Install KVM Guests  
(Control Plane nodes)  
- Bootstrap (temporary node)  
- Control Plane 1  
- Control Plane 2  
- Control Plane 3

Create and Install KVM Guests  
(Worker nodes)  
- Compute Plane 1  
- Compute Plane 2

Verify OCP  
Cluster  
Installation




Pod's IP range

OpenShift SDN


Service's IP range

## install-config.yaml

```
apiVersion: v1
baseDomain: <domain>
compute:
- architecture: s390x
  hyperthreading: Enabled
  name: worker
  replicas: 0
controlPlane:
  architecture: s390x
  hyperthreading: Enabled
  name: master
  replicas: 3
metadata:
  name: <cluster_name>
networking:
  clusterNetwork:
  - cidr: 10.128.0.0/14
    hostPrefix: 23
  networkType: OpenShiftSDN
  serviceNetwork:
  - 172.30.0.0/16
platform:
  none: {}
fips: false
pullSecret: '<pull-secret>'
sshKey: '<ssh-public-key>'
```



# Generate the Ignition Files




1 `# ./openshift-install create manifests --dir=<installation_directory>`  
Modify the `/<installation_directory>/manifests/cluster-scheduler-02-config.yml` `mastersSchedulable` parameter and set its value to `False`

2 `# ./openshift-install create ignition-configs --dir=<installation_directory>`

3 Copy the `bootstrap.ign`, `master.ign` and `worker.ign` to your already pre-configured HTTPd

```
# cp <installation_directory>/*.ign /var/www/html/ignition
# chmod 775 /var/www/html/ignition/*.ign
```


# Prepare the KVM OCP guests



```
1 # wget https://mirror.openshift.com/pub/openshift-v4/s390x/dependencies/rhcos/latest/latest/rhcos-qemu.s390x.qcow2.gz
```

```
2 # dnf install -y gzip
# gunzip rhcos-qemu.s390x.qcow2.gz /var/lib/libvirt/images/
```

# Create Bootstrap



1 # qemu-img create -f qcow2 -F qcow2 -b /var/lib/libvirt/images/rhcos-qemu.s390x.qcow2 /var/lib/libvirt/images/bootstrap.qcow2 120G

2

```
# virt-install --boot kernel=rhcos-kernel,initrd=rhcos-initramfs.img,kernel_args='rd.neednet=1
coreos.inst.install_dev=/dev/vda coreos.live.rootfs_url=http://<bastion_IP>:8080/bin/rhcos-rootfs.img
coreos.inst.ignition_url=http://<bastion_IP>:8080/ignition/bootstrap.ign ip=<bootstrap_IP>::<gateway>:<netmask>:::none
nameserver=<bastion_IP>' --connect qemu:///system --name bootstrap --memory 16384 --vcpus 4 --disk /var/lib/libvirt/
images/bootstrap.qcow2 --accelerate --import --network network=macvtap-net --qemu-commandline="-drive
if=none,id=ignition,format=raw,file=/var/www/html/ignition/bootstrap.ign,readonly=on -device virtio-
blk,serial=ignition,drive=ignition"
```

3 Wait the Bootstrap creation. To verify the install process:

```
# virsh console bootstrap
# journalctl -u bootkube.service
```

## NOTE

Expect many error messages from the bootstrap's log.


Wait for the message: **bootkube.service complete** once the **bootstrap** and **all control plane nodes** are **up** and **running**.

# Red Hat OpenShift Container Platform Install Process

KVM



# Create Masters



```
1 # qemu-img create -f qcow2 -F qcow2 -b /var/lib/libvirt/images/rhcos-qemu.s390x.qcow2 /var/lib/libvirt/images/master1.qcow2 120G
# qemu-img create -f qcow2 -F qcow2 -b /var/lib/libvirt/images/rhcos-qemu.s390x.qcow2 /var/lib/libvirt/images/master2.qcow2 120G
# qemu-img create -f qcow2 -F qcow2 -b /var/lib/libvirt/images/rhcos-qemu.s390x.qcow2 /var/lib/libvirt/images/master3.qcow2 120G
```

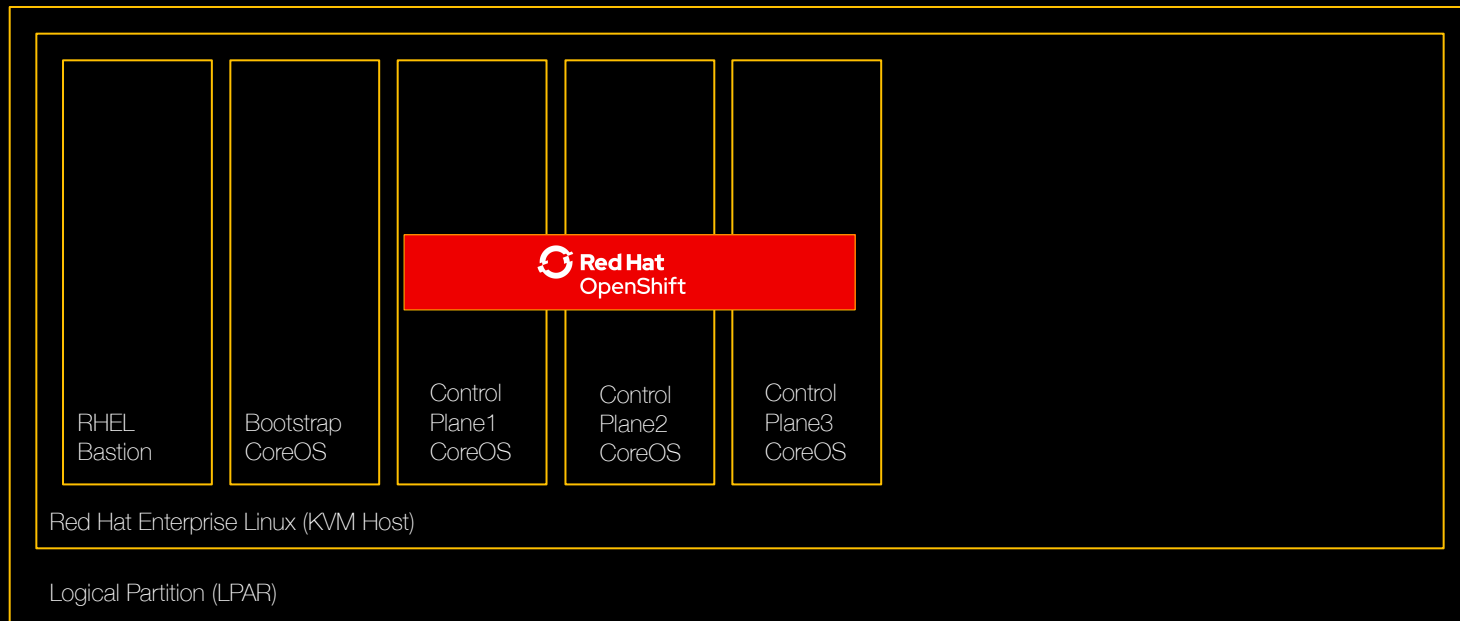
```
2 # virt-install --boot kernel=rhcos-kernel,initrd=rhcos-initramfs.img,kernel_args='rd.neednet=1 coreos.inst.install_dev=/dev/
vda coreos.live.rootfs_url=http://<bastion_IP>:8080/bin/rhcos-rootfs.img coreos.inst.ignition_url=http://<bastion_IP>:8080/
ignition/master.ign ip=<master1_IP>:<gateway>:<netmask>:::none nameserver=<bastion_IP>' --connect qemu:///system --name
master1 --memory 16384 --vcpus 4 --disk /var/lib/libvirt/images/master1.qcow2 --accelerate --import --network
network=macvtap-net --qemu-commandline="-drive if=none,id=ignition,format=raw,file=/var/www/html/ignition/
master.ign,readonly=on -device virtio-blk,serial=ignition,drive=ignition"
```

```
# virt-install --boot kernel=rhcos-kernel,initrd=rhcos-initramfs.img,kernel_args='rd.neednet=1 coreos.inst.install_dev=/dev/
vda coreos.live.rootfs_url=http://<bastion_IP>:8080/bin/rhcos-rootfs.img coreos.inst.ignition_url=http://<bastion_IP>:8080/
ignition/master.ign ip=<master2_IP>:<gateway>:<netmask>:::none nameserver=<bastion_IP>' --connect qemu:///system --name
master2 --memory 16384 --vcpus 4 --disk /var/lib/libvirt/images/master2.qcow2 --accelerate --import --network
network=macvtap-net --qemu-commandline="-drive if=none,id=ignition,format=raw,file=/var/www/html/ignition/
master.ign,readonly=on -device virtio-blk,serial=ignition,drive=ignition"
```

```
# virt-install --boot kernel=rhcos-kernel,initrd=rhcos-initramfs.img,kernel_args='rd.neednet=1 coreos.inst.install_dev=/dev/
vda coreos.live.rootfs_url=http://<bastion_IP>:8080/bin/rhcos-rootfs.img coreos.inst.ignition_url=http://<bastion_IP>:8080/
ignition/master.ign ip=<master3_IP>:<gateway>:<netmask>:::none nameserver=<bastion_IP>' --connect qemu:///system --name
master3 --memory 16384 --vcpus 4 --disk /var/lib/libvirt/images/master3.qcow2 --accelerate --import --network
network=macvtap-net --qemu-commandline="-drive if=none,id=ignition,format=raw,file=/var/www/html/ignition/
master.ign,readonly=on -device virtio-blk,serial=ignition,drive=ignition"
```

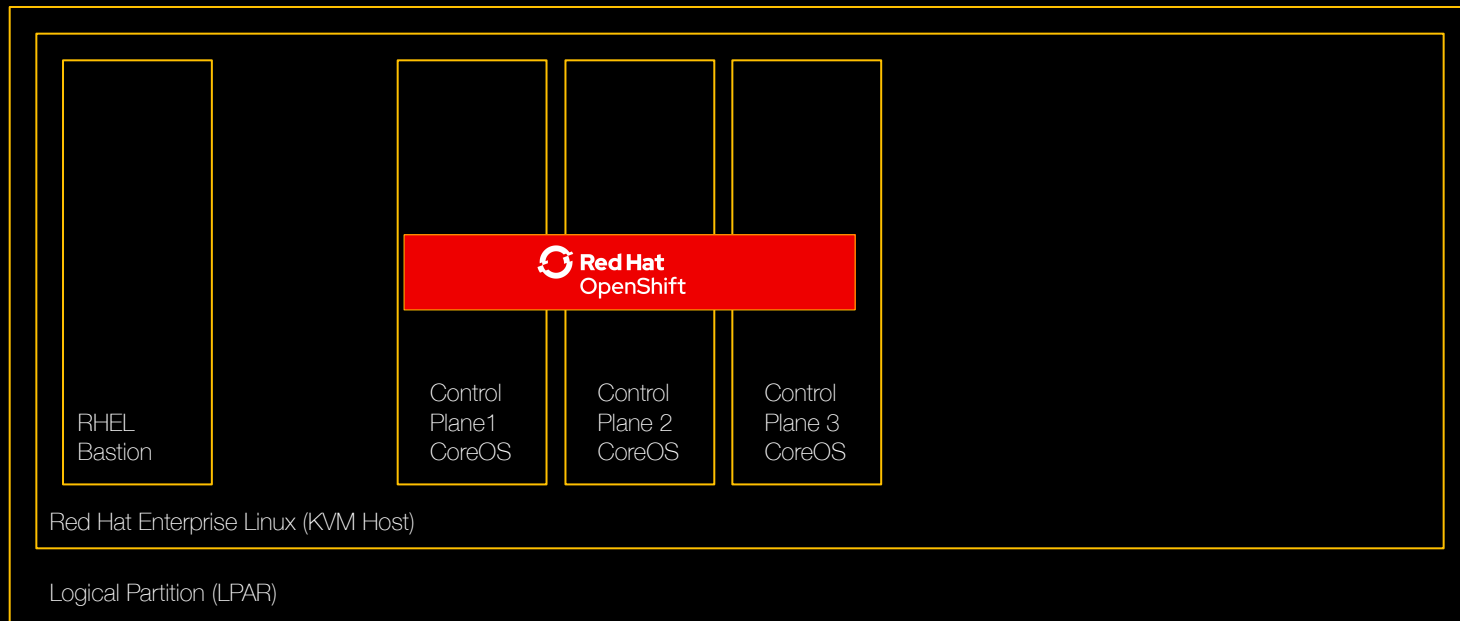
# Red Hat OpenShift Container Platform Install Process

KVM




# Red Hat OpenShift Container Platform Install Process

KVM





# Create Workers



- ```
# qemu-img create -f qcow2 -F qcow2 -b /var/lib/libvirt/images/rhcos-qemu.s390x.qcow2 /var/lib/libvirt/images/worker1.qcow2 120G
# qemu-img create -f qcow2 -F qcow2 -b /var/lib/libvirt/images/rhcos-qemu.s390x.qcow2 /var/lib/libvirt/images/worker2.qcow2 120G
```

- ```
# virt-install --boot kernel=rhcos-kernel,initrd=rhcos-initramfs.img,kernel_args='rd.neednet=1 coreos.inst.install_dev=/dev/vda
coreos.live.rootfs_url=http://<bastion_IP>:8080/bin/rhcos-rootfs.img coreos.inst.ignition_url=http://<bastion_IP>:8080/ignition/
worker.ign ip=<worker1_IP>::<gateway>:<netmask>:::none nameserver=<bastion_IP>' --connect qemu:///system --name worker1 --memory
16384 --vcpus 4 --disk /var/lib/libvirt/images/worker1.qcow2 --accelerate --import --network network=macvtap-net --qemu-
commandline="-drive if=none,id=ignition,format=raw,file=/var/www/html/ignition/worker.ign,readonly=on -device virtio-
blk,serial=ignition,drive=ignition"
```

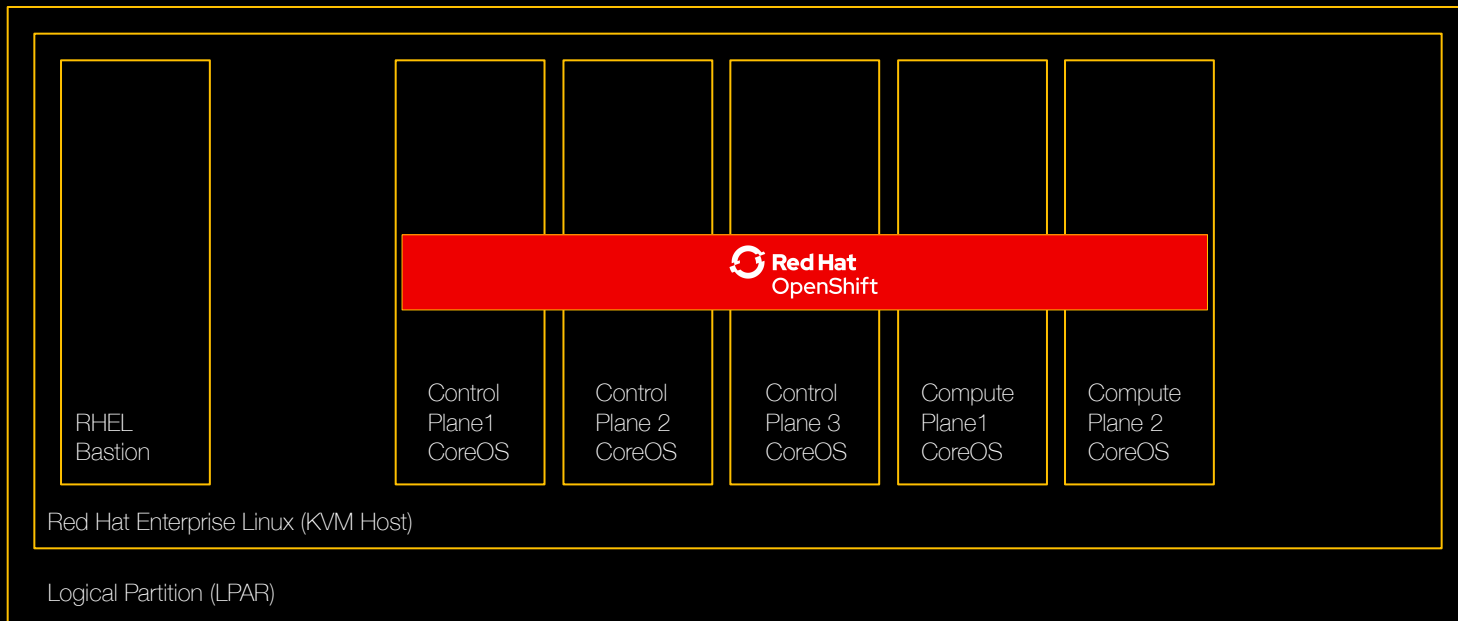
```
# virt-install --boot kernel=rhcos-kernel,initrd=rhcos-initramfs.img,kernel_args='rd.neednet=1 coreos.inst.install_dev=/dev/vda
coreos.live.rootfs_url=http://<bastion_IP>:8080/bin/rhcos-rootfs.img coreos.inst.ignition_url=http://<bastion_IP>:8080/ignition/
worker.ign ip=<worker2_IP>::<gateway>:<netmask>:::none nameserver=<bastion_IP>' --connect qemu:///system --name worker2 --memory
16384 --vcpus 4 --disk /var/lib/libvirt/images/worker2.qcow2 --accelerate --import --network network=macvtap-net --qemu-
commandline="-drive if=none,id=ignition,format=raw,file=/var/www/html/ignition/worker.ign,readonly=on -device virtio-
blk,serial=ignition,drive=ignition"
```

- From your Bastion system, use the OC client to connect the OCP cluster and approve  
Any pending csr certificate:


```
# oc get csr
# oc adm certificate approve <csr-name>
```

# Red Hat OpenShift Container Platform Install Process

KVM



# Verify the Installation



NAME	STATUS	ROLES	AGE	VERSION
master0.<cluster_name>.<domain>	Ready	master	20m	
master1.<cluster_name>.<domain>	Ready	master	20m	
master2.<cluster_name>.<domain>	Ready	master	20m	
worker0.<cluster_name>.<domain>	Ready	worker	20m	
worker1.<cluster_name>.<domain>	Ready	worker	20m	

## 1 Monitor nodes and cluster operators

```
# oc get nodes
# oc get clusteroperators
```

### 1.1 Cluster operator does not come up

```
# oc describe co <clusteroperator>
# oc get pods -n <namespace>
# oc get all -n <namespace>
# oc logs <type>/<name> -n <namespace>
# oc describe nodes
```

NAME	VERSION	AVAILABLE	PROGRESSING	DEGRADED	SINCE
authentication	4.7.7	True	False	False	20m
cloud-credential	4.7.7	True	False	False	20m
cluster-autoscaler	4.7.7	True	False	False	20m
console	4.7.7	True	False	False	5h36m
...					

## 2 Once all operators are marked as State TRUE, conclude the install process:

```
# ./openshift-install --dir=<installation_directory> wait-for install-complete
```

Final Step to complete OCP installation process

## References:

### Red Hat OpenShift Container Platform Reference Architecture

<https://www.ibm.com/docs/en/linux-on-systems?topic=configuration-red-hat-openshift-reference>

### Red Hat OpenShift Container Platform (Static IP)

<https://kvmonz.blogspot.com/2021/03/installing-red-hat-openshift-on-kvm-on-z.html>

### Special Thanks to:

Anderson Augusto Da Silveira

Mainframe Lab Services Field Technical Consultant

[andersonz@br.ibm.com](mailto:andersonz@br.ibm.com)

Thank you  
Grazie  
Merci  
Gracias  
Obrigado  
ありがとう  
谢谢  
Dankeschön



Worldwide zAcceleration

#### Chris Backer

zAcceleration Leader  
[cbacker@us.ibm.com](mailto:cbacker@us.ibm.com)  
zAcceleration  
Red Hat Synergy  
Worldwide IBM Z

#### Filipe Miranda

Cloud Solutions Architect Leader  
[fmiranda@ibm.com](mailto:fmiranda@ibm.com)  
zAcceleration  
Red Hat Synergy  
Worldwide IBM Z

#### Elton de Souza

Chief Architect  
[Elton.desouza@ca.ibm.com](mailto:Elton.desouza@ca.ibm.com)  
Cloud Native Client Success

#### Pat Fruth

Cloud Paks Leader  
[pfruth@us.ibm.com](mailto:pfruth@us.ibm.com)  
zAcceleration  
Red Hat Synergy  
Worldwide IBM Z

#### Roberto Calderon

Cloud z/OS Integration Leader  
[rcalderon@us.ibm.com](mailto:rcalderon@us.ibm.com)  
zAcceleration  
Red Hat Synergy  
Worldwide IBM Z

#### Anna Shugol

Cloud Solutions Engineer  
[anna.shugol@ibm.com](mailto:anna.shugol@ibm.com)  
zAcceleration  
Red Hat Synergy  
Worldwide IBM Z

#### Vic Cross

Cloud Solutions Engineer  
[viccross@au1.ibm.com](mailto:viccross@au1.ibm.com)  
zAcceleration  
Red Hat Synergy  
Worldwide IBM Z